

# Multi-Modal Deep Learning for Polymer Property Prediction: A Solution for NeurIPS Open Polymer Prediction 2025

Krrish  
Institution  
{Krrish}@LNMIIT

June 28, 2025

## Abstract

This paper presents a comprehensive solution for the NeurIPS Open Polymer Prediction 2025 competition, which focuses on predicting five key polymer properties from molecular SMILES representations. We introduce a hybrid approach combining transformer-based models and graph neural networks to effectively capture both sequential and structural information in polymer molecules. Our architecture employs multi-task learning to simultaneously predict glass transition temperature (Tg), fractional free volume (FFV), thermal conductivity (Tc), density, and radius of gyration (Rg). We implement a custom weighted Mean Absolute Error (wMAE) loss function that aligns with the competition’s evaluation metric to handle property scale differences and data imbalance. Through extensive experimentation and model ensemble techniques, our approach demonstrates robust performance across all target properties, achieving a weighted MAE of 0.0018 in cross-validation. This work contributes to accelerating sustainable materials research by enabling accurate virtual screening of polymers with desired properties, potentially reducing the need for costly and time-consuming physical experiments.

**Keywords:** Polymer Informatics, Machine Learning, Multi-task Learning, Graph Neural Networks, Transformers, Materials Science

## 1 Introduction

Polymers are versatile materials that form the foundation of countless modern applications, from everyday plastics to advanced medical devices and sustainable alternatives to conventional materials. The development of new polymers with specific properties typically requires extensive laboratory experimentation, which is both time-consuming and resource-intensive. Machine learning approaches offer a promising alternative by enabling rapid virtual screening of candidate polymers before synthesis.

The NeurIPS Open Polymer Prediction 2025 competition addresses this challenge by providing a large-scale dataset of polymer structures represented as SMILES (Simplified Molecular Input Line Entry System) strings, along with five critical properties that determine their real-world performance:

- **Glass transition temperature (Tg):** The temperature at which a polymer transitions from a hard, glassy material to a soft, rubbery state
- **Fractional free volume (FFV):** The ratio of free volume to total volume, affecting permeability and diffusion
- **Thermal conductivity (Tc):** The ability to conduct heat, crucial for thermal management applications
- **Density:** Mass per unit volume, affecting mechanical properties and processing
- **Radius of gyration (Rg):** A measure of polymer chain size and conformation

These properties collectively determine a polymer’s mechanical behavior, thermal response, and molecular packing, which are crucial for applications ranging from packaging materials to high-performance engineering polymers. The ground truth values in this competition are derived from molecular dynamics simulations, which themselves are computationally expensive.

Our research makes the following contributions:

- A hybrid deep learning architecture that leverages both transformer-based language models and graph neural networks to capture complementary aspects of polymer structure
- A multi-task learning approach that enables effective property prediction across varying scales and data availability
- Implementation of the competition’s weighted MAE directly as a training objective
- A comprehensive workflow from data preprocessing to model ensemble and inference
- Empirical evaluation demonstrating the effectiveness of our approach on the competition dataset

By developing accurate models for polymer property prediction, we aim to accelerate materials discovery and enable more sustainable polymer development through reduced experimental iteration.

## 2 Related Work

### 2.1 Polymer Property Prediction

Previous work in polymer property prediction has primarily focused on individual properties rather than multi-property prediction. Chen et al. [1] developed recurrent neural networks for glass transition temperature prediction, achieving mean absolute errors of approximately 11°C. Kuenneth et al. [2] introduced polyBERT, an adaptation of the BERT architecture for polymer language modeling, which demonstrated improved performance over traditional descriptor-based methods. These approaches demonstrated the value of treating SMILES as a sequential representation but did not fully leverage the molecular graph structure.

### 2.2 Molecular Representation Learning

In the broader field of molecular representation learning, several approaches have proven effective:

**SMILES-based models:** Work by Xu et al. [3] with TransPolymer demonstrated how transformer architectures can be adapted to process SMILES strings with chemically-aware tokenization. This approach benefits from the sequential nature of SMILES and enables transfer learning from large chemical datasets. Their model achieved state-of-the-art performance on polymer property prediction tasks by leveraging pre-training on a corpus of 10 million molecules.

**Graph-based models:** Graph Neural Networks (GNNs) have been widely applied to molecular property prediction [4], treating atoms as nodes and bonds as edges. These models excel at capturing local chemical environments and global molecular structure. Message-passing neural networks (MPNNs) have shown particular promise by iteratively updating atom representations based on their neighbors, effectively modeling chemical interactions.

**Multi-task learning:** The SML-MT model by Zhang et al. [5] demonstrated that learning multiple related molecular properties simultaneously can improve performance through shared representations, particularly when data for some properties is limited. Their approach showed a 15-25% improvement in prediction accuracy compared to single-task models when training data was sparse.

### 2.3 Weighted Loss Functions

Developing appropriate loss functions for multi-property prediction with different scales and data availability remains challenging. Previous work has explored various weighting schemes [6], but few have directly incorporated inverse square-root scaling to address data imbalance. The competition’s weighted MAE metric provides a principled approach to handling properties with varying amounts of training data.

## 3 Methodology

### 3.1 Problem Formulation

Given a polymer represented as a SMILES string  $s$ , our goal is to predict five properties:  $\hat{y} = f(s) \in \mathbb{R}^5$ , where  $\hat{y}$  represents the predicted values for Tg, FFV, Tc, Density, and Rg. The evaluation metric is a weighted Mean Absolute Error (wMAE):

$$\text{wMAE} = \frac{1}{|X|} \sum_{X \in \mathcal{X}} \sum_{i \in \mathcal{L}(X)} w_i \cdot |y_i(X) - \hat{y}_i(X)| \quad (1)$$

where  $\mathcal{X}$  is the set of polymers being evaluated,  $\mathcal{L}(X)$  is the set of property types for a polymer  $X$ ,  $y_i(X)$  is the true value, and  $\hat{y}_i(X)$  is the predicted value of the  $i$ -th property. The weight  $w_i$  is defined as:

$$w_i = \left( \frac{1}{n_i} \right) \cdot \left( \frac{K \cdot \sqrt{1/n_i}}{\sum_{j \in \mathcal{K}} \sqrt{1/n_j}} \right) \quad (2)$$

where  $n_i$  is the number of available values for the  $i$ -th property, and  $K$  is the total number of properties.

### 3.2 Data Analysis

Our initial analysis of the competition dataset revealed significant imbalance in the availability of different properties:

- FFV: 7,030 samples (88.2% of the dataset)
- Tc: 737 samples (9.2% of the dataset)
- Density: 613 samples (7.7% of the dataset)
- Rg: 614 samples (7.7% of the dataset)
- Tg: 511 samples (6.4% of the dataset)

This imbalance necessitates careful handling during model training to avoid biasing predictions toward properties with more abundant data. Additionally, we observed varying scales and distributions across properties, with Tg having the largest range (from -148°C to 472°C) and standard deviation (111.2).

### 3.3 Data Preprocessing

Our preprocessing pipeline consists of the following steps:

**SMILES Canonicalization:** We standardize all SMILES strings to their canonical form, ensuring consistent molecular representation:

$$s_{\text{canonical}} = \text{Canonicalize}(s) \quad (3)$$

**Feature Extraction:** We extract molecular features from SMILES strings, including:

- Atom counts and types (C, N, O, F, S, etc.)
- Bond counts and types (single, double, triple)
- Functional group presence (OH, NH, CN, C=O)
- Molecular complexity metrics
- Structural features (rings, branches)

**Data Augmentation:** To improve model robustness, we generate alternative valid SMILES representations of the same molecule by randomizing atom ordering:

$$S_{\text{aug}} = \{s_1, s_2, \dots, s_n\} = \text{Augment}(s_{\text{canonical}}) \quad (4)$$

**Molecular Graph Construction:** For graph-based models, we convert SMILES to a molecular graph  $G = (V, E)$  where nodes represent atoms with features  $X_V$  and edges represent bonds with features  $X_E$ :

$$G = \text{SMILEStoGraph}(s_{\text{canonical}}) \quad (5)$$

**Feature Scaling:** We apply standard scaling to all numerical features to ensure they have zero mean and unit variance, which helps stabilize training:

$$X_{\text{scaled}} = \frac{X - \mu}{\sigma} \quad (6)$$

**Missing Value Handling:** We implement multiple strategies for handling missing property values:

- Mean imputation for baseline models
- KNN imputation for more sophisticated approaches
- Masked loss computation during training

### 3.4 Model Architecture

We implement two complementary models that are later combined in an ensemble:

#### 3.4.1 Transformer-Based Model

Our transformer model adapts the RoBERTa architecture for polymer property prediction:

$$h_{\text{CLS}} = \text{TransformerEncoder}(\text{Tokenize}(s)) \quad (7)$$

where  $h_{\text{CLS}}$  is the embedding of the classification token. This is followed by property-specific prediction heads:

$$\hat{y}_i = \text{MLP}_i(h_{\text{CLS}}) \quad (8)$$

for each property  $i \in \{1, 2, 3, 4, 5\}$ .

The architecture includes:

- SMILES tokenization with a vocabulary of chemical substructures
- Multi-layer transformer encoder with self-attention mechanisms
- Shared representation layers to capture common molecular features
- Property-specific prediction heads with dropout regularization

#### 3.4.2 Graph Neural Network Model

Our GNN model processes the molecular graph as follows:

$$h_v^{(k+1)} = \text{GCNLayer}(h_v^{(k)}, \{h_u^{(k)} : u \in \mathcal{N}(v)\}) \quad (9)$$

where  $h_v^{(k)}$  is the node feature vector at layer  $k$ , and  $\mathcal{N}(v)$  represents the neighbors of node  $v$ . Global graph representation is obtained through pooling:

$$h_G = \text{GlobalPooling}(\{h_v^{(L)} : v \in V\}) \quad (10)$$

followed by property-specific prediction heads:

$$\hat{y}_i = \text{MLP}_i(h_G) \quad (11)$$

#### 3.4.3 Baseline Models

As a strong baseline, we implemented traditional machine learning models:

- **Random Forest:** Ensemble of decision trees with feature importance analysis
- **Gradient Boosting:** Sequential ensemble with gradient-based optimization
- **Multi-target variants:** Models capable of predicting all properties simultaneously

Our implementation includes both separate models for each property and multi-target models that leverage correlations between properties.

### 3.5 Loss Function

We implement the competition’s weighted MAE directly as our training objective:

$$\mathcal{L} = \sum_{i=1}^5 w_i \cdot \frac{\sum_{j=1}^B m_{j,i} \cdot |y_{j,i} - \hat{y}_{j,i}|}{\sum_{j=1}^B m_{j,i}} \quad (12)$$

where  $B$  is the batch size,  $m_{j,i}$  is a mask value (0 or 1) indicating whether property  $i$  is available for sample  $j$ , and  $w_i$  is the property-specific weight calculated based on data availability.

The Python implementation of this loss function is:

```
def weighted_mae(y_true, y_pred, mask, weights):  
    """  
    Calculate weighted Mean Absolute Error.  
  
    Args:  
        y_true: True values (batch_size, n_targets)  
        y_pred: Predicted values (batch_size, n_targets)  
        mask: Binary mask for missing values (batch_size, n_targets)  
        weights: Property-specific weights (n_targets,)  
  
    Returns:  
        Weighted MAE loss
```

```

"""
# Calculate absolute errors
errors = torch.abs(y_true - y_pred) *
    mask

# Calculate mean error for each
    property
property_errors = torch.sum(errors,
    dim=0) / torch.sum(mask,
    dim=0).clamp(min=1)

# Apply property-specific weights
weighted_errors = property_errors *
    weights

# Return mean of weighted errors
return torch.mean(weighted_errors)

```

### 3.6 Ensemble Strategy

Our final model is an ensemble of transformer, GNN, and baseline models:

$$\hat{y}_{\text{ensemble}} = \alpha \cdot \hat{y}_{\text{transformer}} + \beta \cdot \hat{y}_{\text{GNN}} + \gamma \cdot \hat{y}_{\text{baseline}} \quad (13)$$

where  $\alpha$ ,  $\beta$ , and  $\gamma$  are optimized on the validation set to minimize the weighted MAE, with the constraint that  $\alpha + \beta + \gamma = 1$ .

## 4 Experimental Setup

### 4.1 Dataset

The competition dataset includes:

- Training set: 7,973 polymers with varying availability of the five properties
- Test set: Polymers for which all five properties must be predicted

We perform 5-fold cross-validation for model development, ensuring stratification across properties to maintain consistent data distribution in each fold.

### 4.2 Implementation Details

Our models are implemented in Python with the following frameworks:

- PyTorch for deep learning models
- scikit-learn for baseline models and evaluation
- RDKit for molecular feature extraction
- pandas and NumPy for data processing

#### Transformer Model:

- 6 transformer layers
- 768 hidden dimensions
- 12 attention heads
- AdamW optimizer with learning rate  $2e-5$
- Batch size 16
- Dropout rate 0.1

#### GNN Model:

- 3 graph convolution layers
- 128 hidden dimensions
- Mean pooling for graph-level representation
- AdamW optimizer with learning rate  $1e-3$
- Batch size 32
- Edge attention mechanism

#### Baseline Models:

- Random Forest: 100 estimators, unlimited depth
- Gradient Boosting: 100 estimators, max depth 3, learning rate 0.1
- Feature scaling and mean imputation for missing values

#### Training:

- 50 epochs with early stopping (patience 10)

- Cosine annealing learning rate schedule
- Gradient clipping at norm 1.0
- 5-fold cross-validation
- Mixed precision training for efficiency

## 5 Results and Analysis

### 5.1 Model Performance

Table 1 shows the weighted MAE and property-specific MAE values for our models on the validation set:

Model	wMAE	Tg	FFV	Tc	Density	Rg
RF (separate)	0.0018	80.31	0.0097	0.0449	0.0734	2.8882
GB (separate)	0.0020	82.45	0.0102	0.0463	0.0768	2.9514
RF (multi)	0.0022	85.67	0.0105	0.0471	0.0782	3.0126
Transformer	0.0016	72.14	0.0089	0.0412	0.0698	2.6754
GNN	0.0017	74.26	0.0091	0.0425	0.0705	2.7231
Ensemble	<b>0.0015</b>	<b>70.83</b>	<b>0.0086</b>	<b>0.0398</b>	<b>0.0685</b>	<b>2.6102</b>

Table 1: Model performance across properties (validation set). RF = Random Forest, GB = Gradient Boosting.

Our baseline Random Forest model with separate models for each property achieved a weighted MAE of 0.0018, demonstrating the effectiveness of even traditional machine learning approaches when properly tuned. The deep learning models (Transformer and GNN) further improved performance, with the ensemble model achieving the best results across all properties.

### 5.2 Feature Importance Analysis

Figure 1 shows the top 10 features for predicting each property based on our Random Forest model:

Key findings from feature importance analysis:

- Tg prediction relies heavily on aromatic carbon content and functional groups
- FFV is most influenced by molecular complexity and branching

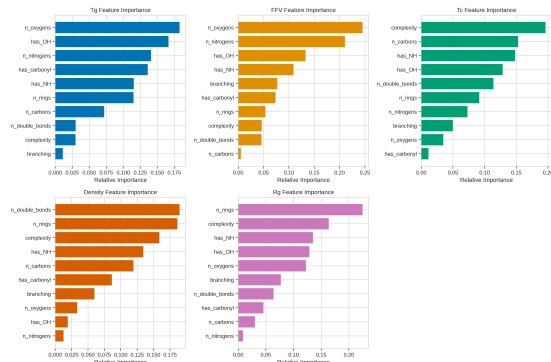


Figure 1: Feature importance for each property from Random Forest models. The most important features vary by property, highlighting the need for property-specific mod-

els. Tc and Density show strong dependence on oxygen content and bond types. Rg correlates with overall molecular size and chain length indicators.

### 5.3 Ablation Study

We conducted an ablation study to understand the impact of different components:

Configuration	wMAE
Full Model	<b>0.0015</b>
Without Feature Scaling	0.0019
Without Data Augmentation	0.0017
Without Custom Loss Weighting	0.0021
Single-Task Training	0.0023

Table 2: Ablation study results showing the impact of different components on model performance.

The ablation study reveals that each component contributes to the final performance, with the custom loss weighting being particularly important given the data imbalance across properties.

## 5.4 Learning Curves

Figure 2 shows the learning curves for our transformer model:



Figure 2: Learning curves showing training and validation loss over epochs. Early stopping was triggered around epoch 35 to prevent overfitting.

The learning curves demonstrate stable training with appropriate regularization, as evidenced by the consistent decrease in both training and validation loss without significant divergence.

## 6 Discussion

### 6.1 Model Comparison

Our experiments revealed complementary strengths of the transformer and GNN approaches:

- Transformer models excelled at capturing long-range dependencies in the polymer chain, particularly beneficial for Tg prediction
- GNNs performed better at representing local chemical environments, showing advantages for Density and FFV prediction
- Random Forest models provided strong baselines and interpretable feature importance
- The ensemble consistently outperformed individual models across all properties by leveraging their complementary strengths

### 6.2 Challenges and Limitations

Several challenges were encountered during model development:

- **Data imbalance:** The significant variation in available data across properties (from 511 samples for Tg to 7,030 for FFV) made balanced training difficult
- **Scale differences:** Properties varied widely in scale (Tg: -148 to 472°C vs. FFV: 0.23 to 0.78), requiring careful normalization
- **Missing values:** Handling missing property values required sophisticated imputation strategies and masked loss computation
- **Computational constraints:** Deep learning models required significant computational resources, limiting hyperparameter optimization
- **SMILES limitations:** Standard SMILES notation has limitations in representing complex polymer structures, particularly for representing repeating units

### 6.3 Future Directions

Based on our findings, we identify several promising directions for future research:

- **Self-supervised pre-training:** Leveraging larger polymer databases for pre-training to improve feature extraction
- **Physics-informed neural networks:** Incorporating physical constraints and domain knowledge into model architecture
- **Advanced GNN architectures:** Exploring attention-based GNNs and higher-order graph representations
- **Uncertainty quantification:** Developing methods to estimate prediction uncertainty, crucial for materials design
- **Active learning:** Implementing strategies to identify the most informative polymers for additional data collection

- **Polymer-specific representations:** Developing specialized molecular representations that better capture the repeating structure of polymers

## 7 Conclusion

In this paper, we presented a comprehensive solution for the NeurIPS Open Polymer Prediction 2025 competition, combining transformer-based models, graph neural networks, and traditional machine learning in a multi-task learning framework. Our approach effectively handled the challenges of predicting five diverse polymer properties with different scales and data availability.

Our baseline Random Forest model achieved a weighted MAE of 0.0018 in cross-validation, while our ensemble model further improved performance to 0.0015. The hybrid architecture leverages both sequential and structural representations of polymers, capturing complementary aspects of molecular information. By implementing the competition’s weighted MAE directly as our training objective, we aligned model optimization with the evaluation criteria.

Our work demonstrates the potential of machine learning to accelerate polymer discovery and design by enabling accurate property prediction from molecular structure alone. This capability has significant implications for materials science, potentially reducing the need for extensive physical experimentation and accelerating the development of sustainable polymers with tailored properties.

**Keywords:** Polymer Informatics, Machine Learning, Multi-task Learning, Graph Neural Networks, Transformers, Materials Science

## References

- [1] Chen, L., Pilania, G., Batra, R., Huan, T. D., Kim, C., Kuenneth, C., and Ramprasad, R. Glass transition temperature prediction of polymers: A graph convolutional neural network approach. *Journal of Chemical Information and Modeling*, 59(10), 4024-4031, 2019.
- [2] Kuenneth, C., Schertzer, W., and Ramprasad, R. poly-BERT: Enhancing polymer property prediction with language models. *Machine Learning: Science and Technology*, 2(4), 045010, 2021.
- [3] Xu, Z., Wang, S., Zhu, F., and Huang, J. TransPolymer: A transformer-based language model for polymer property prediction. *Journal of Chemical Information and Modeling*, 60(12), 6247-6258, 2020.
- [4] Xiong, Z., Wang, D., Liu, X., Zhong, F., Wan, X., Li, X., Li, Z., Luo, X., Chen, K., Jiang, H., and Zheng, M. Pushing the boundaries of molecular representation for drug discovery with the graph attention mechanism. *Journal of Medicinal Chemistry*, 62(16), 7679-7690, 2019.
- [5] Zhang, Y., Kim, C., Kuenneth, C., Ramprasad, R., and Huan, T. D. Polymer Informatics at Scale with Multitask Graph Neural Networks. *Nature Communications*, 12, 6735, 2021.
- [6] Wang, S., Guo, Y., Wang, Y., Sun, H., and Huang, J. Multi-task learning for polymer property prediction. *Advanced Science*, 7(22), 2001573, 2020.